

```

import numpy as np

class Softmax(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the Softmax classifier.
        Note that it has shape (C, D) where C is the number of
        classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims) * 0.0001

    def loss(self, X, y):
        """
        Calculates the softmax loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
            that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """

        # Initialize the loss to zero.
        loss = 0.0
        num_classes = self.W.shape[0]    # C = num_classes
        num_train = X.shape[0]

        exp_a = np.zeros((num_classes, num_train))
        # ===== #
        # YOUR CODE HERE:
        #   Calculate the normalized softmax loss. Store it as the variable loss.
        #   (That is, calculate the sum of the losses of all the training
        #   set margins, and then normalize the loss by the number of
        #   training examples.)
        # ===== #

        for i in np.arange(num_train):

            Loss = 0.0

            class_scores = np.dot(self.W, X[i, :].T)    # calculating class scores (C x 1 vector)
            class_scores -= np.max(class_scores)         # considering the possible issue for num

            exp_a[:, i] = np.exp(class_scores)           # turning class scores to probabilities

            Loss -= np.log(exp_a[y[i], i] / np.sum(exp_a[:, i]))

            # p[:, i] = exp_a[:, i] / np.sum(exp_a[:, i])    # p now is a valid probability
            # print(p[:, i])

            loss += Loss
            # print(Loss, i)

        pass
        loss /= num_train
        # ===== #
        # END YOUR CODE HERE
        # ===== #

        return loss

    def loss_and_grad(self, X, y):
        """
        Same as self.loss(X, y), except that it also returns the gradient.

```

Output: grad — a matrix of the same dimensions as W containing
the gradient of the loss with respect to W.

"""

Initialize the loss and gradient to zero.

loss = 0.0

grad = np.zeros_like(self.W)

grad_tmp = np.zeros_like(self.W)

num_classes = self.W.shape[0] # C = num_classes

num_train = X.shape[0]

=====

YOUR CODE HERE:

Calculate the softmax loss and the gradient. Store the gradient

as the variable grad.

=====

exp_a = np.zeros((num_classes, num_train))

for i in np.arange(num_train):

Loss = 0.0

class_scores = np.dot(self.W, X[i, :].T)

calculating class scores (C x 1 vector)

class_scores -= np.max(class_scores)

considering the possible issue for num

exp_a[:, i] = np.exp(class_scores)

turning class scores to probabilities

Loss -= np.log(exp_a[y[i], i] / np.sum(exp_a[:, i]))

#if i==0:

grada = np.zeros(X.shape[1])

for j in range(num_classes):

if j != y[i]:

grad_tmp[j, :] = X[i, :].T * (exp_a[j, i] / np.sum(exp_a[:, i]))

else:

grad_tmp[j, :] = X[i, :].T * (exp_a[j, i] / np.sum(exp_a[:, i])) - X[i, :].T

grad += grad_tmp

loss += Loss

pass

loss /= num_train

grad /= num_train

=====

END YOUR CODE HERE

=====

return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):

"""

sample a few random elements and only return numerical

in these dimensions.

"""

for i in np.arange(num_checks):

ix = tuple([np.random.randint(m) for m in self.W.shape])

oldval = self.W[ix]

self.W[ix] = oldval + h # increment by h

fxph = self.loss(X, y)

self.W[ix] = oldval - h # decrement by h

fxmh = self.loss(X, y) # evaluate f(x - h)

self.W[ix] = oldval # reset

grad_numerical = (fxph - fxmh) / (2 * h)

grad_analytic = your_grad[ix]

rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) + abs(grad_analytic))

print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical, grad_analytic, rel_error))

def fast_loss_and_grad(self, X, y):

"""

```

A vectorized implementation of loss_and_grad. It shares the same
    inputs and ouptuts as loss_and_grad.
"""
loss = 0.0
grad = np.zeros(self.W.shape) # initialize the gradient as zero

# ===== #
# YOUR CODE HERE:
#   Calculate the softmax loss and gradient WITHOUT any for loops.
# ===== #

num_train = X.shape[0]
num_classes = self.W.shape[0]

#   # vectorized loss calculation #
class_scores_matrix = np.dot(self.W,X.T) # calculating class scores matrix (C x m):
class_scores_matrix -= np.max(class_scores_matrix) # considering the possible issue
exp_a = np.exp(class_scores_matrix) # calculating the exponents

#   y_exp = np.array(exp_a[y, np.arange(0, class_scores_matrix.shape[1])])
#   #print(exp_a[:, :3])
#   #print(y[:3])
#   #print(y_exp[:3])

#   tt = np.sum(exp_a, axis=0)
#   tt2 = np.divide(tt, y_exp)
#   print(num_train)
#   tt3 = np.power(tt2, 1/num_train)
#   loss = np.log(np.prod(tt3))

(C, D) = self.W.shape
N = X.shape[0]

scores = np.dot(self.W, X.T)
scores -= np.max(scores) # shift by log C to avoid numerical instability

y_mat = np.zeros(shape = (C, N))
y_mat[y, range(N)] = 1

# matrix of all zeros except for a single wx + log C value in each column that corresponds
# quantity we need to subtract from each row of scores
correct_wx = np.multiply(y_mat, scores)

# create a single row of the correct wx_y + log C values for each data point
sums_wy = np.sum(correct_wx, axis=0) # sum over each column

exp_scores = np.exp(scores)
sums_exp = np.sum(exp_scores, axis=0) # sum over each column
result = np.log(sums_exp)

result -= sums_wy

loss = np.sum(result)
loss /= num_train

# vectorized gradient calculation #
exp_a_sum = np.sum(exp_a, axis=0)

y_mat_corres = np.zeros(shape = (num_classes, num_train))
y_mat_corres[y, range(num_train)] = 1
sum_exp_scores = np.sum(exp_a, axis=0)
sum_exp_scores = 1.0 / exp_a_sum # division by sum over columns
exp_a *= sum_exp_scores
grad = np.dot(exp_a, X)
grad -= np.dot(y_mat_corres, X)
grad /= num_train

# ===== #
# END YOUR CODE HERE
# ===== #

```

```

    return loss , grad

def train(self , X, y, learning_rate=1e-3, num_iters=100,
        batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there are N
        training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i] = c
        means that X[i] has label 0 <= c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each step.
    - verbose: (boolean) If true, print progress during optimization.

    Outputs:
    A list containing the value of the loss function at each training iteration.
    """
    num_train, dim = X.shape
    num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes

    self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of self.W

    # Run stochastic gradient descent to optimize W
    loss_history = []

    for it in np.arange(num_iters):
        X_batch = None
        y_batch = None

        # ===== #
        # YOUR CODE HERE:
        # Sample batch_size elements from the training data for use in
        # gradient descent. After sampling,
        # - X_batch should have shape: (dim, batch_size)
        # - y_batch should have shape: (batch_size,)
        # The indices should be randomly generated to reduce correlations
        # in the dataset. Use np.random.choice. It's okay to sample with
        # replacement.
        # ===== #
        mask = np.random.choice(num_train, batch_size, replace=True)

        X_batch = X[mask] # (dim, batch_size)
        y_batch = y[mask] # (batch_size,)

        pass

        # ===== #
        # END YOUR CODE HERE
        # ===== #

        # evaluate loss and gradient
        loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
        loss_history.append(loss)

        # ===== #
        # YOUR CODE HERE:
        # Update the parameters, self.W, with a gradient step
        # ===== #
        pass

        self.W = self.W - learning_rate*grad

        # ===== #
        # END YOUR CODE HERE
        # ===== #

        if verbose and it % 100 == 0:
            print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

    return loss_history

def predict(self, X):

```

```

"""
Inputs:
- X: N x D array of training data. Each row is a D-dimensional point.

Returns:
- y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
  array of length N, and each element is an integer giving the predicted
  class.
"""

y_pred = np.zeros(X.shape[1])
# ===== #
# YOUR CODE HERE:
#   Predict the labels given the training data.
# ===== #

y_pred = np.argmax(np.exp(self.W.dot(X.T)), axis=0)

pass
# ===== #
# END YOUR CODE HERE
# ===== #

return y_pred

```