

```

import numpy as np
import pdb

"""
This code was based off of code from cs231n at Stanford University, and modified for ece239as
"""
class SVM(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the SVM. Note that it has shape (C, D)
        where C is the number of classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims)

    def loss(self, X, y):
        """
        Calculates the SVM loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
            that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """

        # compute the loss and the gradient
        num_classes = self.W.shape[0]
        num_train = X.shape[0]
        loss = 0.0

        for i in np.arange(num_train):
            lost = 0.0
            # ===== #
            # YOUR CODE HERE:
            # Calculate the normalized SVM loss, and store it as 'loss'.
            # (That is, calculate the sum of the losses of all the training
            # set margins, and then normalize the loss by the number of
            # training examples.)
            # ===== #
            y_est = np.dot(self.W, X[i, :].T)
            class_est = y_est[y[i]]
            #print(y_est)
            #print(class_est)
            #print(y[i])
            for j in range(num_classes):
                if j != y[i]:
                    #print(class_est)
                    #print(y_est[j])
                    lost += np.maximum(0, 1 - class_est + y_est[j])
            loss += lost

        pass

        #print(lost)
        loss /= num_train
        print(lost)
        print(num_train)

        # ===== #
        # END YOUR CODE HERE
        # ===== #

        return loss

    def loss_and_grad(self, X, y):

```

```
"""
```

```
    Same as self.loss(X, y), except that it also returns the gradient.
```

```
    Output: grad — a matrix of the same dimensions as W containing
            the gradient of the loss with respect to W.
```

```
"""
```

```
# compute the loss and the gradient
num_classes = self.W.shape[0]
num_train = X.shape[0]
loss = 0.0
grad = np.zeros_like(self.W)      # for whole data
grad_tmp = np.zeros_like(self.W)  # for each data

for i in np.arange(num_train):
# ===== #
# YOUR CODE HERE:
#   Calculate the SVM loss and the gradient. Store the gradient in
#   the variable grad.
# ===== #
    lost = 0.0
    y_est = np.dot(self.W, X[i, :].T)
    class_est = y_est[y[i]]

    #print(y_est)
    #print(class_est)
    #print(y[i])

    #if i==0:
    grada = np.zeros(X.shape[1])
#lost: ith training data contribution to cost function
    for j in range(num_classes):
        if j != y[i]:
            #print(class_est)
            #print(y_est[j])
            lost += np.maximum(0, 1 - class_est + y_est[j])
            if (1 - class_est + y_est[j] > 0):
                grad_tmp[j, :] = X[i, :].T
                grada -= X[i, :].T
            else:
                grad_tmp[j, :] = np.zeros(X.shape[1])

        grad_tmp[y[i], :] = grada

    grad += grad_tmp
    loss += lost

loss /= num_train
grad /= num_train

#print(grad[0, :])
pass

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical
    in these dimensions.
    """

    for i in np.arange(num_checks):
        ix = tuple([np.random.randint(m) for m in self.W.shape])

        oldval = self.W[ix]
        self.W[ix] = oldval + h # increment by h
        fxph = self.loss(X, y)
        self.W[ix] = oldval - h # decrement by h
```

```

    fxmh = self.loss(X,y) # evaluate f(x - h)
    self.W[ix] = oldval # reset

    grad_numerical = (fxph - fxmh) / (2 * h)
    grad_analytic = your_grad[ix]
    rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) + abs(grad_analytic))
    print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical, grad_analytic, rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and outputs as loss_and_grad.
    """
    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero

    # ===== #
    # YOUR CODE HERE:
    # Calculate the SVM loss WITHOUT any for loops.
    # ===== #

    scores = np.dot(self.W, X.T)

    bias_by_one = np.ones(scores.shape) # adding bias 1 in vector form

    training_scores = np.ones(scores.shape) * [scores[y, np.arange(0, scores.shape[1])]]

    Loss = scores + bias_by_one - training_scores
    Loss_mod = Loss
    # performing max function against zero in vector form
    # also we should not count the remaining that are equal to 1
    Loss_mod[Loss_mod < 0] = 0
    Loss_mod[y, np.arange(0, scores.shape[1])] = 0 # not counting elements that are equal to 1
    loss = np.sum(Loss_mod)

    # Averaging over number of training data
    num_train = X.shape[0]
    loss /= num_train

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    # ===== #
    # YOUR CODE HERE:
    # Calculate the SVM grad WITHOUT any for loops.
    # ===== #

    Loss_mod2 = Loss

    Loss_mod2[Loss_mod2 < 0] = 0 # we don't care about the ones that have negative margins
    Loss_mod2[Loss_mod2 > 0] = 1 # the positive margins contribute to the loss with x then
    Loss_mod2[y, np.arange(0, scores.shape[1])] = 0 # we take care of these rows corresponding

    Loss_mod2[y, np.arange(0, scores.shape[1])] = -1 * np.sum(Loss_mod2, axis=0) # rows corresponding
    grad = np.dot(Loss_mod2, X)

    # Averaging over number of training data
    num_train = X.shape[0]
    grad /= num_train

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
          batch_size=200, verbose=False):

```

"""

Train this linear classifier using stochastic gradient descent.

Inputs:

- X: A numpy array of shape (N, D) containing training data; there are N training samples each of dimension D.
- y: A numpy array of shape (N,) containing training labels; $y[i] = c$ means that $X[i]$ has label $0 \leq c < C$ for C classes.
- learning_rate: (float) learning rate for optimization.
- num_iters: (integer) number of steps to take when optimizing
- batch_size: (integer) number of training examples to use at each step.
- verbose: (boolean) If true, print progress during optimization.

Outputs:

A list containing the value of the loss function at each training iteration.

"""

```
num_train, dim = X.shape
```

```
num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes
```

```
self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of self.W
```

```
# Run stochastic gradient descent to optimize W
```

```
loss_history = []
```

```
for it in np.arange(num_iters):
```

```
    X_batch = None
```

```
    y_batch = None
```

```
    # ===== #
```

```
    # YOUR CODE HERE:
```

```
    # Sample batch_size elements from the training data for use in
```

```
    # gradient descent. After sampling,
```

```
    # - X_batch should have shape: (dim, batch_size)
```

```
    # - y_batch should have shape: (batch_size,)
```

```
    # The indices should be randomly generated to reduce correlations
```

```
    # in the dataset. Use np.random.choice. It's okay to sample with
```

```
    # replacement.
```

```
    # ===== #
```

```
    mask = np.random.choice(num_train, batch_size, replace=True)
```

```
    X_batch = X[mask] # (dim, batch_size)
```

```
    y_batch = y[mask] # (batch_size,)
```

```
    # ===== #
```

```
    # END YOUR CODE HERE
```

```
    # ===== #
```

```
    # evaluate loss and gradient
```

```
    loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
```

```
    loss_history.append(loss)
```

```
    # ===== #
```

```
    # YOUR CODE HERE:
```

```
    # Update the parameters, self.W, with a gradient step
```

```
    # ===== #
```

```
    self.W = self.W - learning_rate*grad
```

```
    # ===== #
```

```
    # END YOUR CODE HERE
```

```
    # ===== #
```

```
    if verbose and it % 100 == 0:
```

```
        print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
```

```
return loss_history
```

```
def predict(self, X):
```

```
    """
```

Inputs:

- X: N x D array of training data. Each row is a D-dimensional point.

Returns:

- y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional array of length N, and each element is an integer giving the predicted

```

class .
"""
y_pred = np.zeros(X.shape[1])

# ===== #
# YOUR CODE HERE:
#   Predict the labels given the training data with the parameter self.W.
# ===== #

y_pred = np.argmax(self.W.dot(X.T), axis=0)

# ===== #
# END YOUR CODE HERE
# ===== #

return y_pred

```